

PAR Laboratory Assignment

Lab 4: Divide and Conquer parallelism with OpenMP: Sorting

Your names here, group, date, and include boada usernames

Add delivery date here



UNIVERSITAT POLITÈCNICA
DE CATALUNYA
BARCELONATECH

How to:

Document contents

The length of the document is expected to be maximum 8 pages (including figures and tables) with a font size of at least 11pt.

You have to add comments and observations following the questions that you have found along the document indicating **For the deliverable:** in each section (in this laboratory we have included those **For the deliverable** questions in the notebook with sentences in gray under each **Comments/Observations**). DON'T include code in the document except for the fragment modified with respect to the original one: just provide a code except if you consider necessary.

Finally, as you know, this course contributes to the **transversal competence "Tercera llengua"**. Deliver your material in English if you want this competence to be evaluated. Please refer to the "Rubrics for the third language competence evaluation" document to know the *Rubric* that will be used.

Submission

Only PDF format for this document will be accepted. Deliverable assignment will be opened in *Atenea* and set the appropriate dates for the delivery. You also have to deliver the complete C source codes for Tareador instrumentation and all the OpenMP parallelisation strategies that you have done. Your lab professor should be able to re-execute the parallel codes based on the files you deliver. You will have to **deliver TWO files**, one with the document in PDF format and one compressed file (**tgz**, **.gz** or **.zip**) with the requested C source codes.

1

Laboratory 4 notebook

1.1 Task decomposition analysis with *Tareador*

1.1.1 Leaf and Tree Analysis with *Tareador*

This part refers to subsection 2.2 of the practice document.

Include the TDG of both *leaf* and *tree* strategies *Tareador* analysis.

Comments/Observations

Identify the points of the code where you should include synchronizations to guarantee the dependences for each task decomposition strategy. Which directives/annotations/clauses will you use to guarantee those dependences in the *OpenMP* implementations?

1.2 Leaf and Tree strategies in *OpenMP*

1.2.1 Analysis with *modelfactors*

This part refers to point 4 of subsection 3.1 and subsection 3.2 of the practice document.

Include the three tables generated for the *leaf* and *tree* strategies. Neither *leaf* nor *tree* strategies are showing a good parallel performance.

Comments/Observations

Which of the factors do you think is making the parallelisation efficiency so low in the case of the *leaf*? Which of the following parameters do you think is making the *tree* parallelisation better than the *leaf* version? Several options to think about: parallel fraction, in-execution efficiency (related with the overheads of sync and sched reported in the third table), number of tasks and their execution time, load balancing, ... Is the granularity of both (*leaf* and *tree*) strategies influencing the parallel performance (see overheads in *modelfactor* tables for both strategies)? What is the number of tasks created (see *modelfactor* tables for both strategies)?

1.2.2 Analysis with *Paraver*

This part refers to point 5 of subsection 3.1 and subsection 3.2 of the practice document.

Include parts (zoom in) of the *Paraver* visualisations that help you explain the lack of scalability. In particular, and for both strategies, we think it would be good to show those parts that show examples of:

1) the amount of task executed in parallel, 2) which threads are executing tasks and 3) which thread/s is/are creating tasks?.

Comments/Observations

Is the program generating enough tasks to simultaneously feed all threads? How many? How many threads are creating tasks?

1.3 Task granularity control: the *cut-off* mechanism

1.3.1 Analysis with *modelfactors*

This part refers to points 4 and 5 of subsection 3.3 of the practice document.

Include the tables of *modelfactors* for cut-off level equal to 4 and the information of the number of tasks generated for each of the cut-off levels used.

Comments/Observations

Which is the best value for cut-off? Does it significantly change with the number of threads used?

1.3.2 Analysis with *Paraver*

This part refers to point 5 of subsection 3.3 of the practice document.

Include the *Paraver* trace visualizations you have analyzed for the execution with 8 threads.

Comments/Observations

Is the instantaneous parallelism achieved larger than one along all the execution trace?

1.3.3 Granularity Analysis

This part refers to point 3 of subsection 3.2 and 7 of subsection 3.3 of the practice document.

Include the postscript file generated by `submit-cutoff-omp.sh` to explore different cut-off levels for 8 threads.

Comments/Observations

Which is the best value for cut-off for this problem size (look at the script to figure out the size of the problem) and 8 threads?

Include the strong scalability plot generated for the best cut-off level.

Comments/Observations

Are the speedup for the complete and multisort function close to the speed-up ideal? Reason the difference in speedup based on *Paraver* trace you have analyzed and the parallelization of your `multisort-omp.c` code.

2

Shared-memory parallelisation with *OpenMP* task using dependencies

This part refers to section 4 of the practice document.

Include:

- Tables of *modelfactors*.
- Strong Scalability plot.
- The *Paraver* traces analyzed.

Comments/Observations

Is this parallel implementation better or worse compare to *OpenMP* versions of previous chapter in terms of performance? In terms of programmability, was this new version simpler to code?